
Buffer Overflow Attack: Attack Techniques

The Problem

```
void foo(char *s) {  
    char buf[10];  
    strcpy(buf,s);  
    printf("buf is %s\n",s);  
}
```

...

```
foo("thisstringistolongforfoo");
```

Exploitation

- ❑ The general idea is to give servers very large strings that will overflow a buffer.
- ❑ For a server with sloppy code – it's easy to crash the server by overflowing a buffer.
- ❑ It's sometimes possible to actually make the server do whatever you want (instead of crashing).

Necessary Background

- ❑ C functions and the stack.
- ❑ A little knowledge of assembly/machine language.
- ❑ How system calls are made (at the level of machine code level).
- ❑ **exec ()** system calls
 - ❑ How to “guess” some key parameters.

What is a Buffer Overflow?

❑ **Intent**

- ❑ Arbitrary code execution
 - ❑ Spawn a remote shell or infect with worm/virus
- ❑ Denial of service
 - ❑ Cause software to crash
 - ❑ E.g., ping of death attack

❑ **Steps**

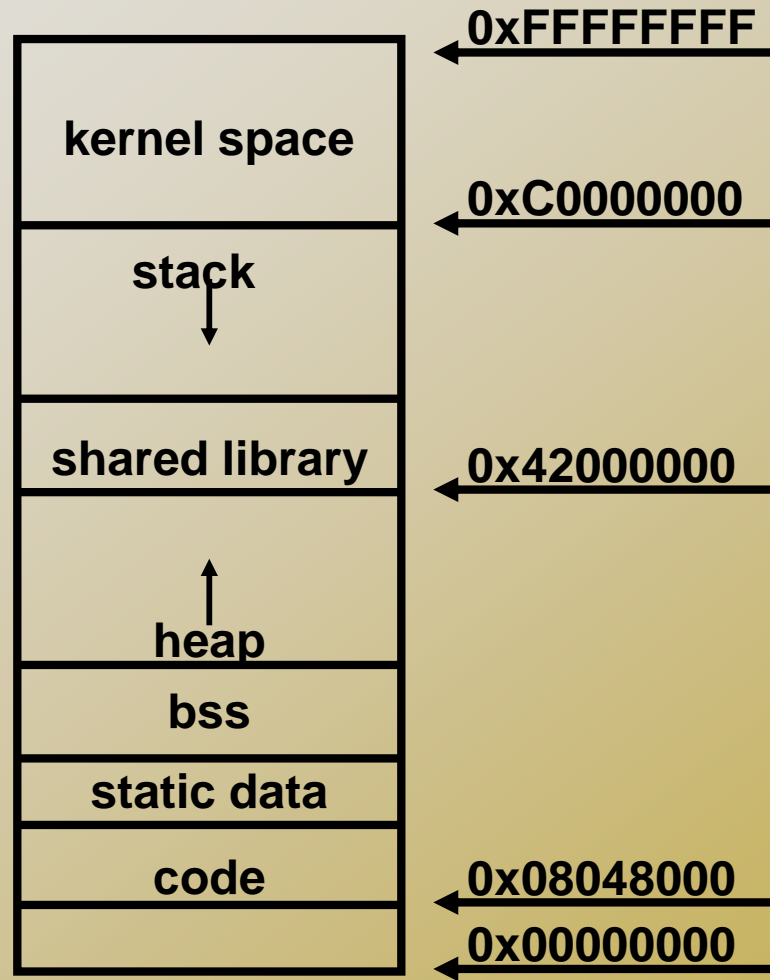
- ❑ Inject attack code into buffer
- ❑ Overflow return address
- ❑ Redirect control flow to attack code
- ❑ Execute attack code

Attack Possibilities

- ❑ **Targets**
 - ❑ Stack, heap, static area
 - ❑ Parameter modification (non-pointer data)
 - ❑ Change parameters for existing call to exec()
 - ❑ Change privilege control variable
- ❑ **Injected code vs. existing code**
- ❑ **Absolute vs. relative address dependence**
- ❑ **Related Attacks**
 - ❑ Integer overflows
 - ❑ Format-string attacks

❑ Stack Overflow Overview

Address Space

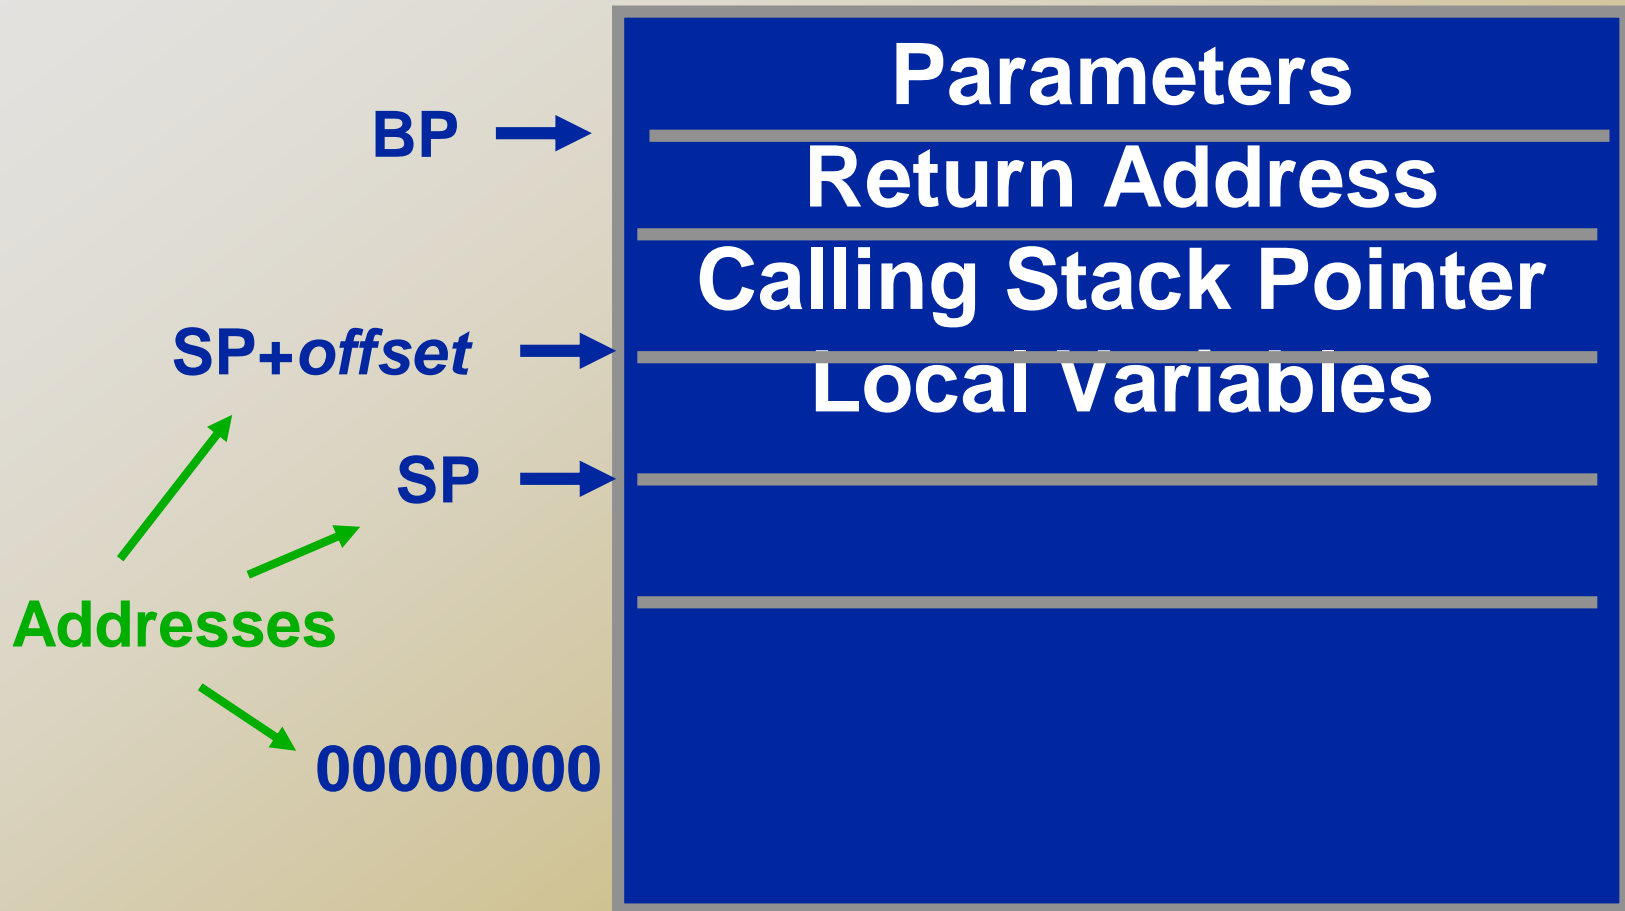


From Dawn Song's RISE: <http://research.microsoft.com/projects/SWSecInstitute/slides/Song.ppt>

C Call Stack

- ❑ **C Call Stack**
 - ❑ When a function call is made, the return address is put on the stack.
 - ❑ Often the values of parameters are put on the stack.
 - ❑ Usually the function saves the stack frame pointer (on the stack).
 - ❑ Local variables are on the stack.

A Stack Frame



SP: stack pointer BP: base/frame pointer

Sample Stack

18

addressof(y=3) return address
saved stack pointer

buf

y

x

```
x=2;  
foo(18);  
y=3;
```

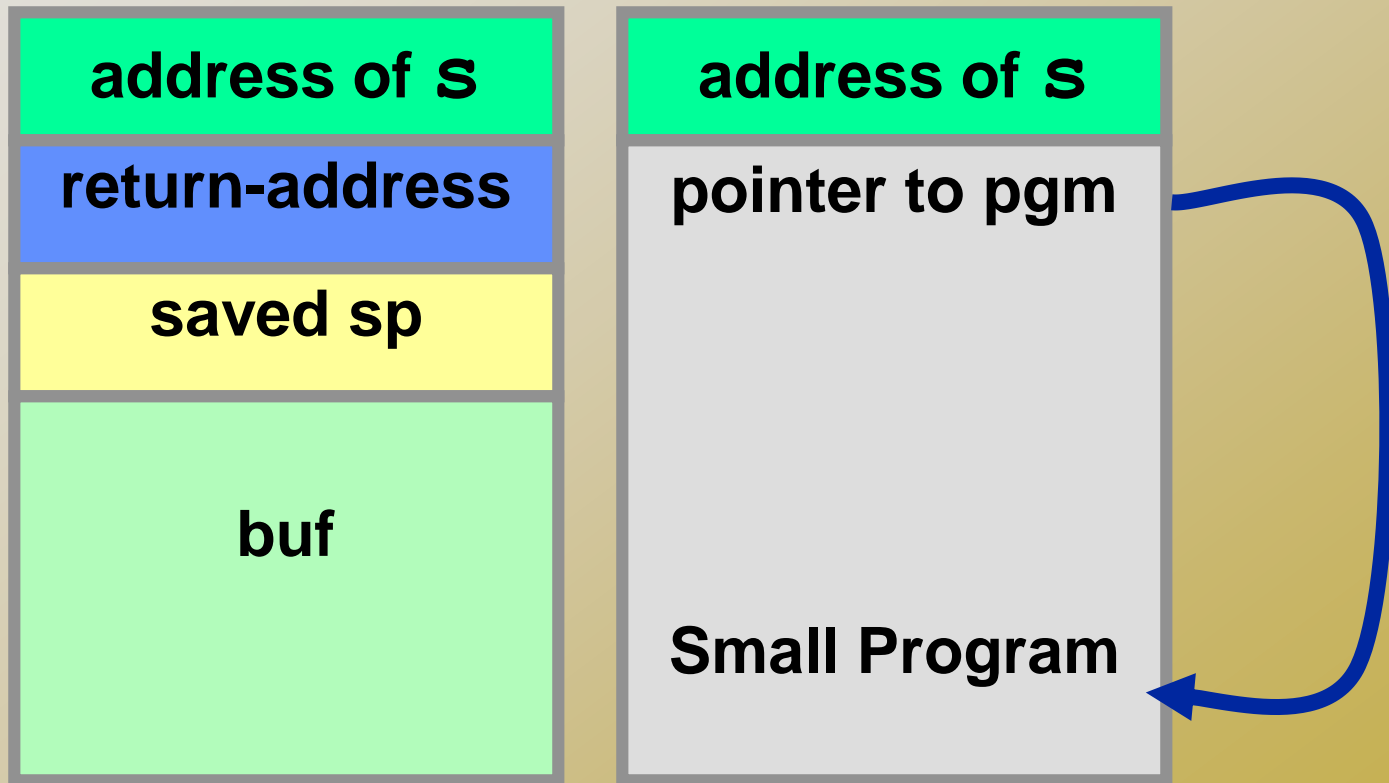
```
void foo(int j) {  
    int x,y;  
    char buf[100];  
    x=j;  
    ...  
}
```

“Smashing the Stack”*

- ❑ The general idea is to overflow a buffer so that it overwrites the return address.
- ❑ When the function is done it will jump to whatever address is on the stack.
- ❑ We put some code in the buffer and set the return address to point to it!

Before and After

```
void foo(char *s) {  
    char buf[100];  
    strcpy(buf, s);  
    ...  
}
```



❑ What causes buffer overflow?

Example: gets()

```
char buf[20];  
gets(buf); // read user input until  
           // first EoL or EOF character
```

- ❑ Never use **gets**
- ❑ Use **fgets(buf, size, stdout)** instead

Example: strcpy()

```
char dest[20];  
strcpy(dest, src); // copies string src to  
dest
```


- ❑ **strcpy** assumes **dest** is long enough ,
and assumes **src** is null-terminated
- ❑ Use **strncpy(dest, src, size)** instead

Spot the defect! (1)

```
char buf[20];  
char prefix[] = "http://";  
...  
strcpy(buf, prefix);  
    // copies the string prefix to buf  
strncat(buf, path, sizeof(buf));  
    // concatenates path to the string buf
```

Spot the defect! (1)

```
char buf[20];  
char prefix[] = "http://";  
...  
strcpy(buf, prefix);  
    // copies the string prefix to buf  
strncat(buf, path, sizeof(buf));  
    // concatenates path to the string buf
```



strncat's 3rd parameter is number of
chars to copy, not the buffer size

Another common mistake is giving **sizeof(path)** as 3rd argument...

Spot the defect! (2)

```
char src[9];  
char dest[9];
```

```
char base_url = "www.ru.nl";  
strncpy(src, base_url, 9);  
    // copies base_url to src  
strcpy(dest, src);  
    // copies src to dest
```

base_url is 10 chars long, incl. its null terminator, so **src** won't be null-terminated

so **strcpy** will overrun the buffer **dest**

Example: **strcpy** and **strncpy**

- ❑ Don't replace **strcpy(dest, src)** by
 - ❑ **strncpy(dest, src, sizeof(dest))**
- ❑ but by
 - ❑ **strncpy(dest, src, sizeof(dest)-1)**
 - ❑ **dest[sizeof(dest)-1] = '\0';**
 - ❑ if **dest** should be null-terminated!
- ❑ A strongly typed programming language could of course enforce that strings are always null-terminated...

Spot the defect! (3)

```
char *buf;  
int i, len;  
read(fd, &len, sizeof(len));  
buf = malloc(len);  
read(fd, buf, len);
```

Spot the defect! (3)

```
char *buf;  
int i, len;  
read(fd, &len, sizeof(len));  
buf = malloc(len);  
read(fd, buf, len);
```

Didn't check if negative



len cast to unsigned and negative length overflows



- ❑ **Memcpy() prototype:**
 - ❑ `void *memcpy(void *dest, const void *src, size_t n);`
- ❑ **Definition of size_t: typedef unsigned int size_t;**

Implicit Casting Bug

- ❑ **A signed/unsigned or an implicit casting bug**
 - ❑ **Very nasty – hard to spot**
- ❑ **C compiler doesn't warn about type mismatch between signed int and unsigned int**
 - ❑ **Silently inserts an implicit cast**

Spot the defect! (4)

```
char *buf;  
int i, len;  
read(fd, &len, sizeof(len));  
if (len < 0)  
    {error ("negative length"); return; }  
buf = malloc(len+5);  
read(fd,buf,len);  
buf[len] = '\0'; // null terminate buf
```

May results in integer overflow



Spot the defect! (5)

```
#define MAX_BUF = 256
```

```
void BadCode (char* input)
```

```
{    short len;
```

```
    char buf[MAX_BUF];
```

```
    len = strlen(input);
```

```
    if (len < MAX_BUF) strcpy(buf,input);
```

```
}
```

What if input is longer than 32K ?


len will be a negative number,

due to integer overflow

hence: potential buffer overflow

Spot the defect! (6)

```
char buff1[MAX_SIZE], buff2[MAX_SIZE];
// make sure it's a valid URL and will fit
if (! isValid(url)) return;
if (strlen(url) > MAX_SIZE - 1) return;
// copy url up to first separator, ie. first '/', to buff1
out = buff1;
do {
    // skip spaces
    if (*url != ' ') *out++ = *url;
} while (*url++ != '/');
strcpy(buff2, buff1);
...
```



what if there is no '/' in the URL?

Loop termination (exploited by Blaster worm)

Spot the defect! (7)

```
#include <stdio.h>
int main(int argc, char* argv[])
{ if (argc > 1)
printf(argv[1]);
return 0;
}
```

This program is vulnerable to **format string** attacks, where calling the program with strings containing special characters can result in a buffer overflow attack.

Format String Attacks

- ❑ `int printf(const char *format [, argument]...);`
 - ❑ `snprintf, wsnprintf ...`
- ❑ What may happen if we execute **`printf(string);`**
 - ❑ Where **`string`** is user-supplied ?
 - ❑ If it contains special characters, eg `%s`, `%x`, `%n`, `%hn`?

Format String Attacks

- ❑ **Why this could happen?**
 - ❑ Many programs delay output message for batch display:
 - ❑ **fprintf(STDOUT, err_msg);**
 - ❑ Where the err_msg is composed based on user inputs
 - ❑ If a user can change err_msg freely, format string attack is possible

Format String Attacks

- ❑ **%x** reads and prints 4 bytes from stack
 - ❑ this may leak sensitive data
- ❑ **%n** writes the number of characters printed so far onto the stack
 - ❑ this allow stack overflow attacks...
- ❑ C format strings break the “don’t mix data & code” principle.
- ❑ “Easy” to spot & fix:
 - ❑ replace **printf(str)** by **printf(“%s”, str)**

Use Unix Machine in Department

- ❑ The Unix machine:
cseaimlcsbs.iem.edu.in
- ❑ Must use SSH to connect
 - ❑ Find free SSH clients on Internet
 - ❑ E.g., Putty (command line based)
 - ❑ http://en.wikipedia.org/wiki/Ssh_client
 - ❑ Find a GUI-based SSH client
- ❑ Username:
- ❑ Default password:
the first initial of your last name in
uppercase and the last 5 digits of your
PID

Example of “%x” --- Memory leaking

```
#include <stdio.h>
void main(int argc, char **argv){
    int a1=1; int a2=2;
    int a3=3; int a4=4;
    printf(argv[1]);
}
```

```
czou@:~$ ./test
czou@eustis:~$ ./test "what is this?"
what is this?czou@eustis:~$
czou@eustis:~$
czou@eustis:~$ ./test "%x %x %x %x %x %x"
4 3 2 1 bfc994b0 bfc99508czou@eustis:~$
czou@eustis:~$
```

Bfc994b0: saved stack pointer

Bfc99508: return address

```
#include <stdio.h>
void foo(char *format){
    int a1=11; int a2=12;
    int a3=13; int a4=14;
    printf(format);
}
void main(int argc, char **argv){
    foo(argv[1]);
    printf("\n");
}
$./format-x-subfun "%x %x %x %x : %x,  %x,  %x  "
80495bc e d c : b,  bffff7e8,  80483f4
```



Four variables



Return address

-
- ❑ **What does this string (“%x:%x:%s”) do?**
 - ❑ Prints first two words of stack memory
 - ❑ Treats next stack memory word as memory addr and prints everything until first '\0'
 - ❑ Could segment fault if goes to other program's memory

-
- ❑ **Use obscure format specifier (%n) to write any value to any address in the victim's memory**
 - ❑ %n --- write 4 bytes at once
 - ❑ %hn --- write 2 bytes at once
 - ❑ **Enables attackers to mount malicious code injection attacks**
 - ❑ Introduce code anywhere into victim's memory
 - ❑ Use format string bug to overwrite return address on stack (or a function pointer) with pointer to malicious code

Example of “%n”---- write data in memory

```
#include <stdio.h>
void main(int argc, char **argv){
    int bytes;
    printf(“%s%n\n”, argv[1], &bytes);
    printf(“You input %d characters\n”, bytes);
}
```

czou@eustis:~\$./test hello

hello

You input 5 characters

Function Pointer Overwritten

- ❑ Function pointers: (used in attack on PHP 4.0.2)



- ❑ Overflowing buf will override function pointer.
- ❑ Harder to defend than return-address overflow attacks

Test by Yourself

```
#include <stdio.h>
void main(void){
    /* short x = 32767;*/
    unsigned short x = 65535;
    x = x + 1;
    printf("x= %d\n", x);
}
```

```
#include <stdio.h>
int main()
{
    /* short x = 32767;*/ signed
    short x = -65535;
    x = x + 1;
    printf("x= %d\n", x);
    return 0;
}
```

- Try to run it to see how overflow happens.
- Modify the x definition to see other integer overflow cases

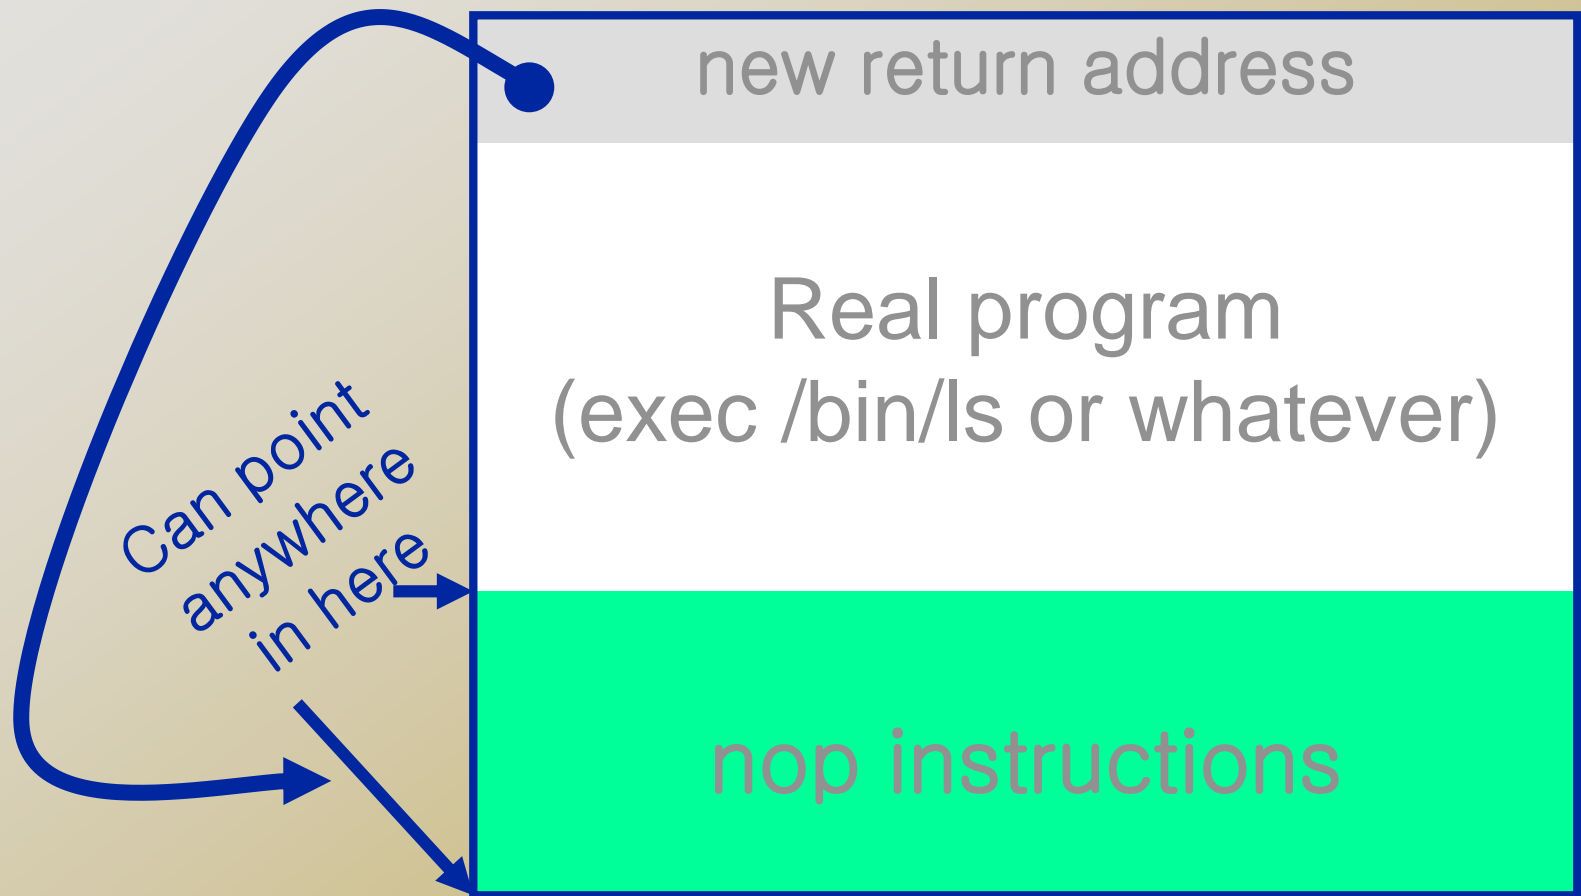
Two Techniques for Generating Stack Overflow Codes

They make attacking relatively easier

NOPs

- ❑ Most CPUs have a *No-Operation* instruction – it does nothing but advance the instruction pointer.
- ❑ Usually we can put a bunch of these ahead of our program (in the string).
- ❑ As long as the new return-address points to a NOP we are OK.


Using NOPs



Estimating the stack size

- ❑ We can also guess at the location of the return address relative to the overflowed buffer.
- ❑ Put in a bunch of new return addresses!

Estimating the Location



Attack Code in Stack Overflow

- ❑ **Most Stack Overflow attacks use “shell code” as the attacking code**
 - ❑ Shell code is small (< 100 bytes)
- ❑ **Attacker will obtain a shell (remote login) for arbitrary command execution**
 - ❑ Remote login without password or account
- ❑ **Shell generated is mostly root privileged (has the same privilege as the compromised program)**
- ❑ **Additional attacking codes can be downloaded from the shell created**
 - ❑ Most Internet worms use this way

Linux Shell Code

Testshell.c:

```
static char shellcode[] =  
"\xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07\x89\x46\x0c\xb0\x0b\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80\x31\xdb\x89\xd8\x40\xcd\x80\xe8\xdc\xff\xff\xff/bin/sh";  
  
int main(void)  
{  
    void (*code)() = (void *)shellcode;  
    //printf("Shellcode length: %d\n", strlen(shellcode));  
    code();  
    return 0;  
}
```